

CSCI U511 Operating Systems

Fall 2019

Division of Mathematics and Computer Science

USC Upstate

Homework-2

Weights: 50pts

Due Date: September 17, 2019 (Tuesday).

Purpose: From doing this assignment, you will learn how to implement an OS command interpreter (shell); and how to use some system calls (such as fork() and exec()).

Turn in Method: Turn in your program via Blackboard by 5pm of September 17, 2019.

Important: You must put your name in the first line of your program as program comments. Deduction will be given for failing to provide this information at the beginning of your program! Deduction (or no credit) will be given if you copy any code from online.

Description: Building a shell

UNIX shells:

The OS command interpreter is the program that people interact with in order to launch and control programs. On UNIX systems, the command interpreter is often called **shell**: a user-level program that gives people a command-line interface to launching, suspending, and killing other programs. sh, ksh, csh, tcsh, bash, ... are all examples of UNIX shells. It might be useful to look at the manual pages of these shells, for example, type "man csh".

The most rudimentary shell is structured as the following loop:

1. Print out a prompt (e.g., "CSCI1/511Shell\$ ");
2. Read a line from the user;
3. Parse the line into the program name and an array of parameters;
4. Use the fork() system call to spawn a new child process;
 - o The child process then uses the exec() system call (or one of its variants) to launch the specified program;
 - o The parent process (the shell) uses the wait() system call (or one of its variants) to wait for the child to terminate;
5. Once the child (the launched program) finishes, the shell repeats the loop by jumping to 1.

Although most commands people type on the shell prompt are the names of other UNIX programs (such as ps or cat), shells also recognize some special commands (called internal commands) that are not program names. For example, the exit command terminates the shell, and the cd command changes the current working directory. Shells directly make system calls to execute these commands, instead of forking a child process to handle them.

Requirements in detail:

Your job is to implement a very primitive shell that knows how to launch new programs in the foreground and the background. It should also recognize a few internal commands.

CSCI U511 Operating Systems

Fall 2019

Division of Mathematics and Computer Science

USC Upstate

More specifically, it should support the following features.

- It should recognize the internal commands: `exit`, `jobs`, and `cd`. `exit` should use the `exit()` system call to terminate the shell. `cd` uses the `chdir()` system call to change to a new directory.
- If the command line does not indicate any internal commands, it should be in the following form:
 `<program name> <arg1> <arg2> <argN> [&]`
Your shell should invoke the program, passing it the list of arguments in the command line. The shell must wait until the started program completes unless the user runs it in the background (with `&`).

To allow users to pass arguments you need to parse the input line into words separated by whitespace (spaces and `\t` tab characters). You might try to use `strtok_r()` for parsing (check the manual page of `strtok_r()` and Google it for examples of using it). In case you wonder, `strtok_r()` is a user-level utility, not a system call. This means this function is fulfilled without the help of the operating system kernel. To make the parsing easy for you, you can assume the `'&'` token (when used) is separated from the last argument with one or more spaces or `\t` tab characters.

The shell runs programs using two core system calls: `fork()` and `execvp()`. In short, `fork()` creates an exact copy of the currently running process, and is used by the shell to spawn a new process. The `execvp()` call is used to overload the currently running program with a new program, which is how the shell turns a forked process into the program it wants to run. In addition, the shell must wait until the previously started program completes unless the user runs it in the background (with `&`). This is done with the `wait()` system call or one of its variants (such as `waitpid()`). All these system calls can fail due to unforeseen reasons. You should check their return status and report errors if they occur.

No input the user gives should cause the shell to exit (except when the user types `exit` or `Ctrl+D`). This means your shell should handle errors gracefully, no matter where they occur. Even if an error occurs in the middle of a long pipeline, it should be reported accurately and your shell should recover gracefully. In addition, your shell should not generate leaking open file descriptors. **Hint:** you can monitor the current open file descriptors of the shell process through the `/proc` file system.