


Graphs II

Digraphs, Strongly Connective
Component, Topological Sorting,
and Minimum Spanning Tree


Digraphs

- A digraph is a graph whose edges are all directed
 - Short for “directed graph”
- Applications
 - one-way streets
 - flights
 - task scheduling




Digraph Properties

- A graph $G=(V,E)$ such that
 - Each edge goes in one direction:
 - Edge (a,b) goes from a to b , but not b to a
- If G is simple, $m \leq n \cdot (n - 1)$
- If we keep in-edges and out-edges in separate adjacency lists, we can perform listing of incoming edges and outgoing edges in time proportional to their size




Digraph Application

- Scheduling: edge (a,b) means task a must be completed before b can be started





Directed DFS

- We can specialize the traversal algorithms (DFS and BFS) to digraphs by traversing edges only along their direction
- In the directed DFS algorithm, we have four types of edges
 - discovery edges
 - back edges
 - forward edges
 - cross edges
- A directed DFS starting at a vertex s determines the vertices reachable from s




Reachability

- DFS tree rooted at v: vertices reachable from v via directed paths





Strong Connectivity

- Each vertex can reach all other vertices




Strong Connectivity Algorithm

- Pick a vertex v in G
- Perform a DFS from v in G
 - If there's a w not visited, print "no"
- Let G' be G with edges reversed
- Perform a DFS from v in G'
 - If there's a w not visited, print "no"
 - Else, print "yes"
- Running time: $O(n+m)$




Strongly Connected Components

- Maximal subgraphs such that each vertex can reach all other vertices in the subgraph
- Can also be done in $O(n+m)$ time using DFS, but is more complicated (similar to biconnectivity).




Transitive Closure

- Given a digraph G , the transitive closure of G is the digraph G^* such that
 - G^* has the same vertices as G
 - if G has a directed path from u to v ($u \neq v$), G^* has a directed edge from u to v
- The transitive closure provides reachability information about a digraph



Computing the Transitive Closure

- We can perform DFS starting at each vertex
 - $O(n(n+m))$




If there's a way to get from A to B and from B to C, then there's a way to get from A to C.

Alternatively ... Use dynamic programming: **The Floyd-Warshall Algorithm**

Floyd-Warshall Transitive Closure

- Idea #1: Number the vertices $1, 2, \dots, n$.
- Idea #2: Consider paths that use only vertices numbered $1, 2, \dots, k$, as intermediate vertices:

Uses only vertices numbered $1, \dots, k$
(add this edge if it's not already in)



Floyd-Warshall's Algorithm

- Number vertices v_1, \dots, v_n
- Compute digraphs G_0, \dots, G_n
 - $G_0 = G$
 - G_k has directed edge (v_i, v_j) if G has a directed path from v_i to v_j with intermediate vertices in $\{v_1, \dots, v_k\}$
- We have that $G_n = G^*$
- In phase k , digraph G_k is computed from G_{k-1}
- Running time: $O(n^3)$, assuming areAdjacent is $O(1)$ (e.g., adjacency matrix)

Algorithm *FloydWarshall*(G)

Input digraph G

Output transitive closure G^* of G

$i \leftarrow 1$

for all $v \in G.\text{vertices}()$

denote v as v_i

$i \leftarrow i + 1$

$G_0 \leftarrow G$

for $k \leftarrow 1$ **to** n **do**

$G_k \leftarrow G_{k-1}$

for $i \leftarrow 1$ **to** n ($i \neq k$) **do**

for $j \leftarrow 1$ **to** n ($j \neq i, k$) **do**

if $G_{k-1}.\text{areAdjacent}(v_i, v_k) \wedge$

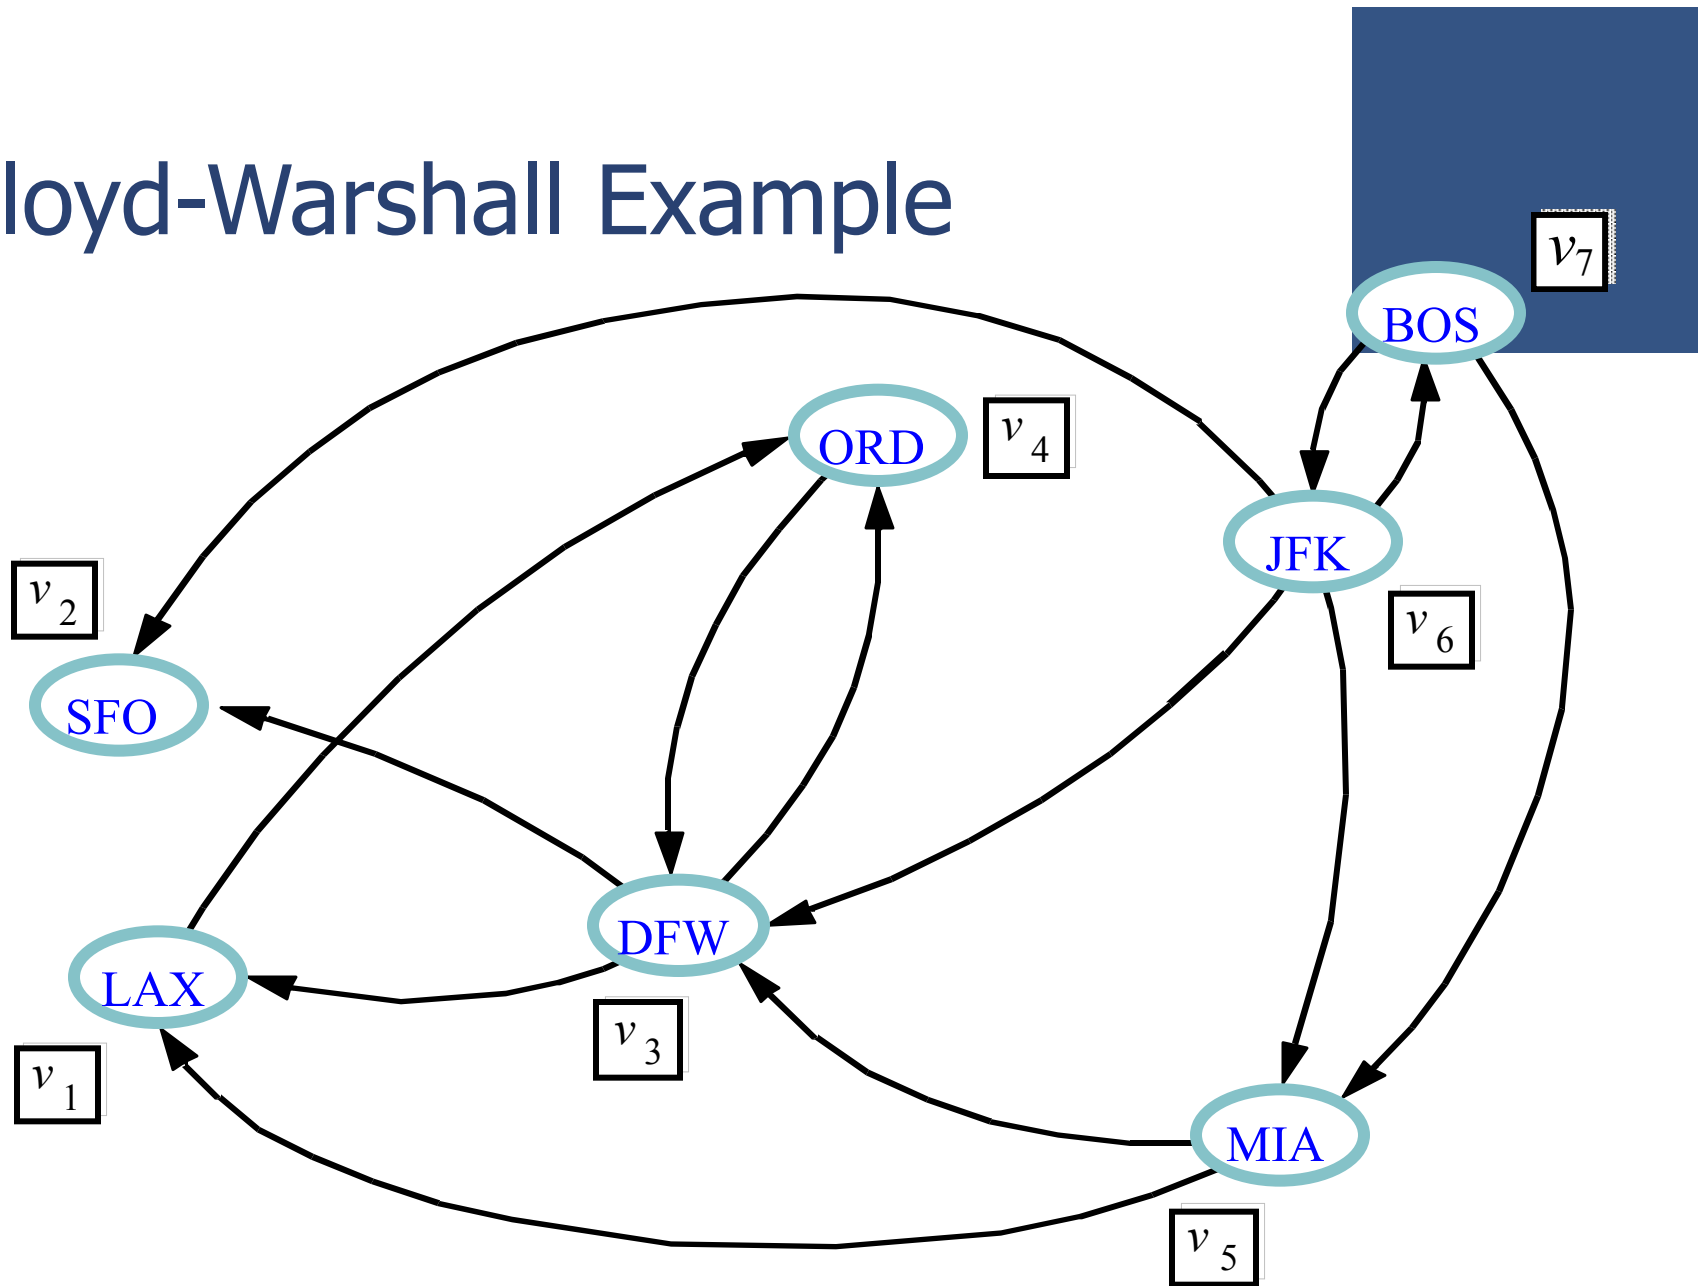
$G_{k-1}.\text{areAdjacent}(v_k, v_j)$

if $\neg G_k.\text{areAdjacent}(v_i, v_j)$


$G_k.\text{insertDirectedEdge}(v_i, v_j, k)$

return G_n

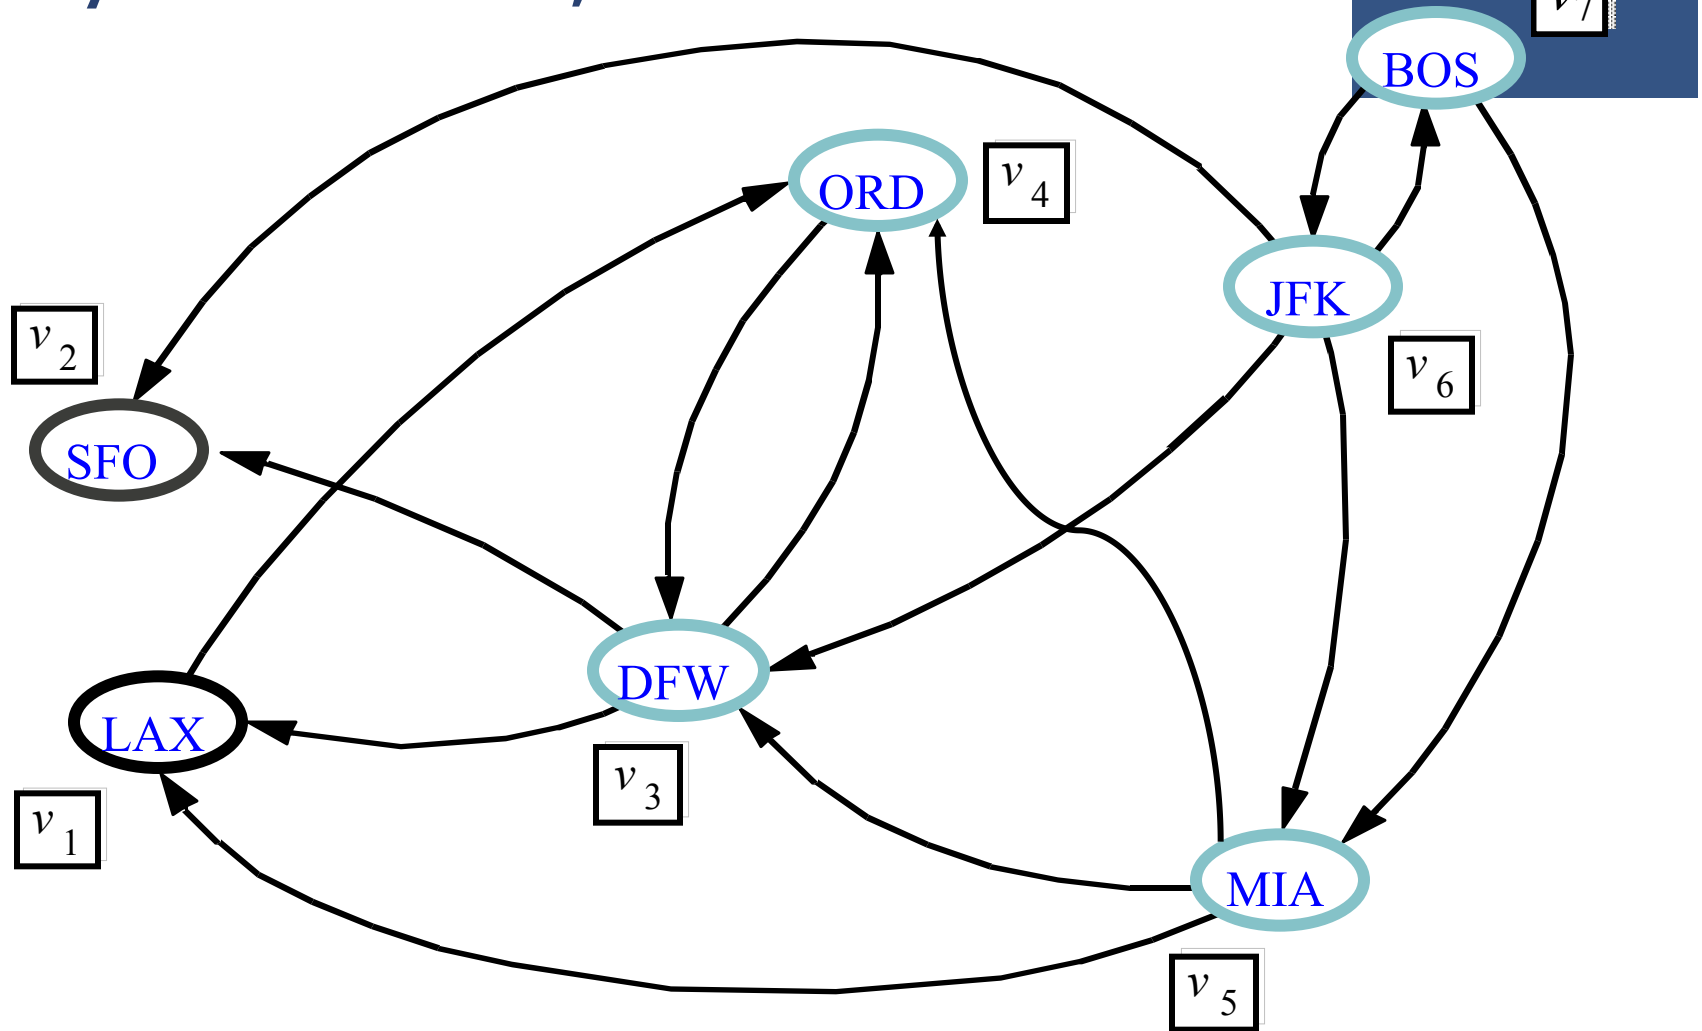
Floyd-Warshall Example



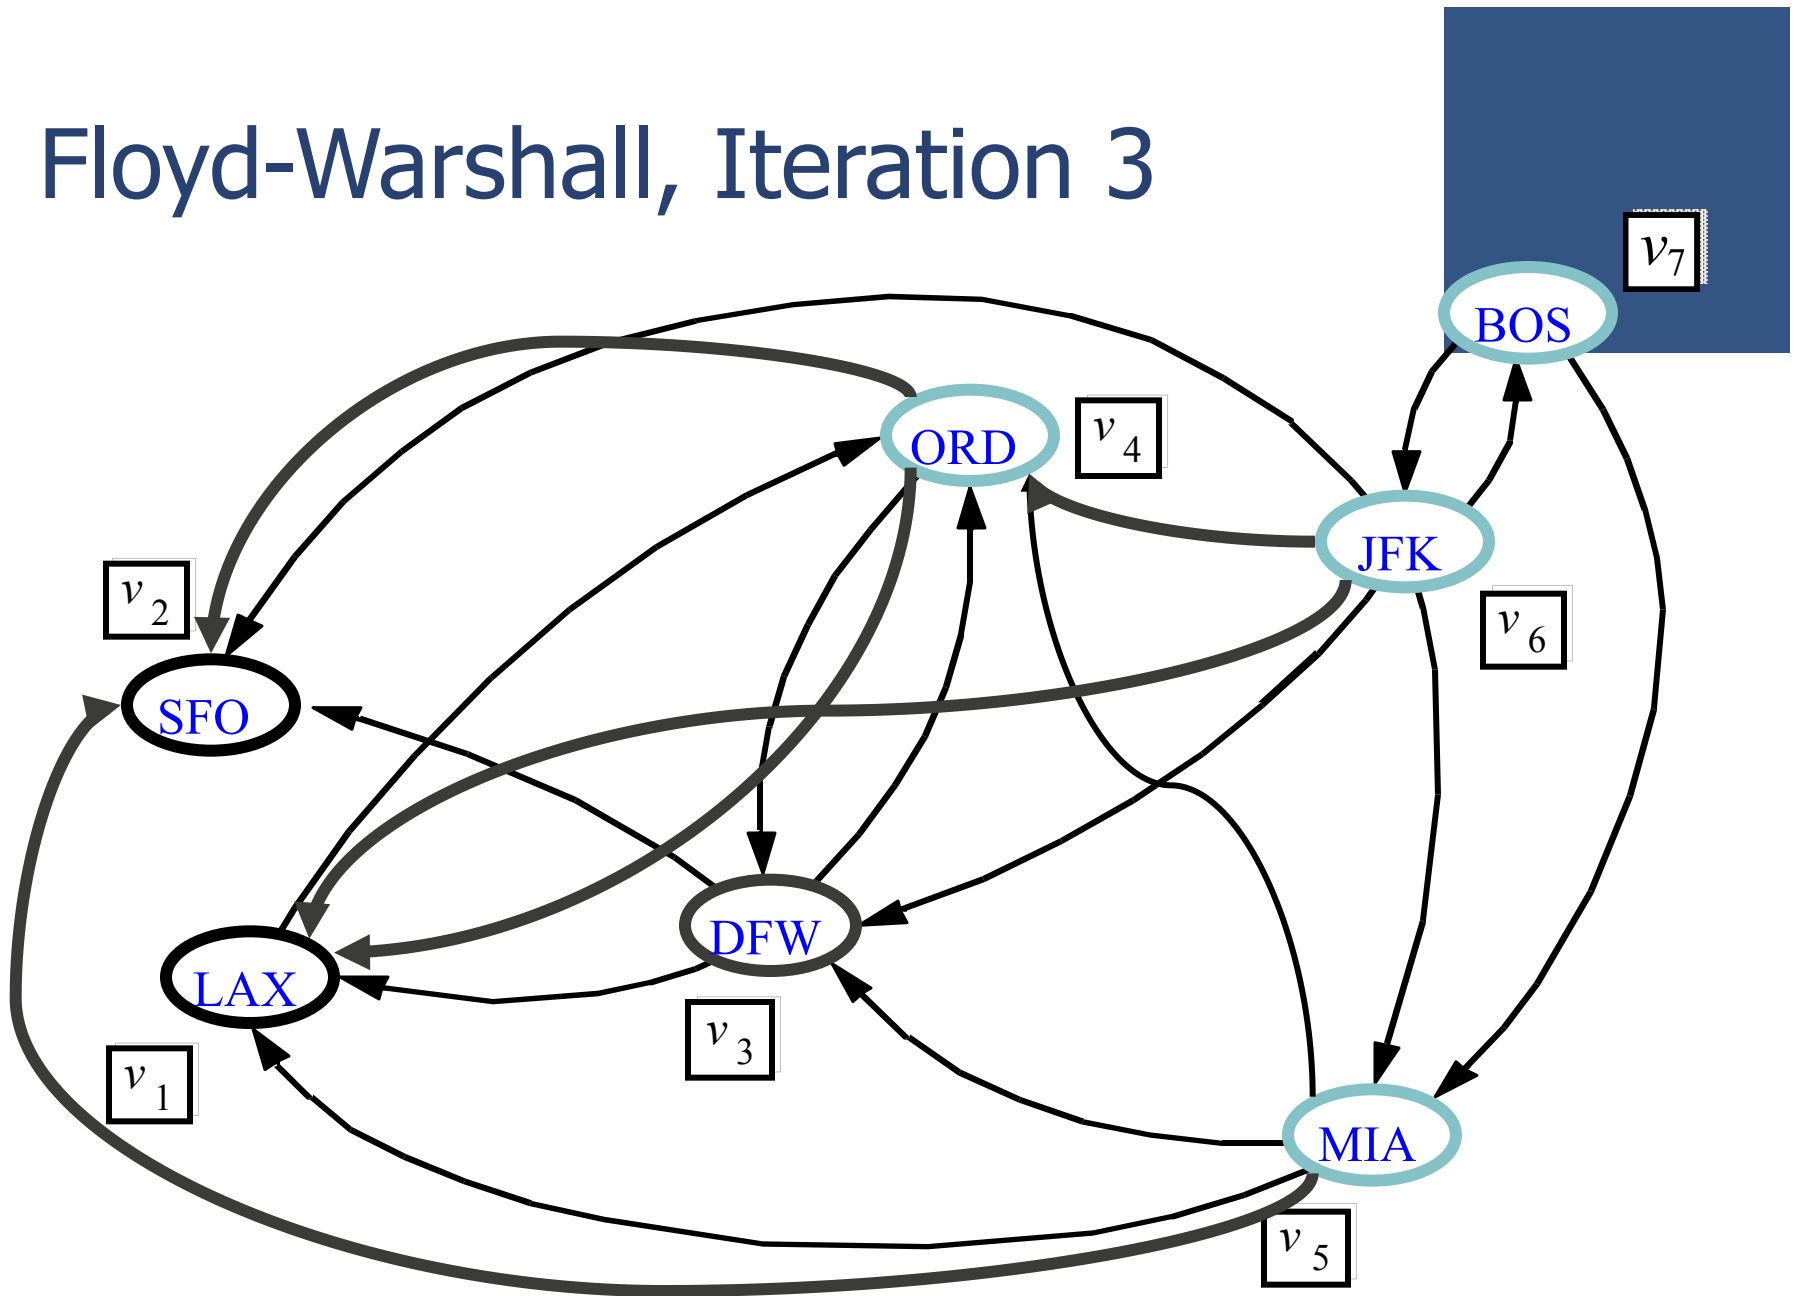
Floyd-Warshall, Iteration 1



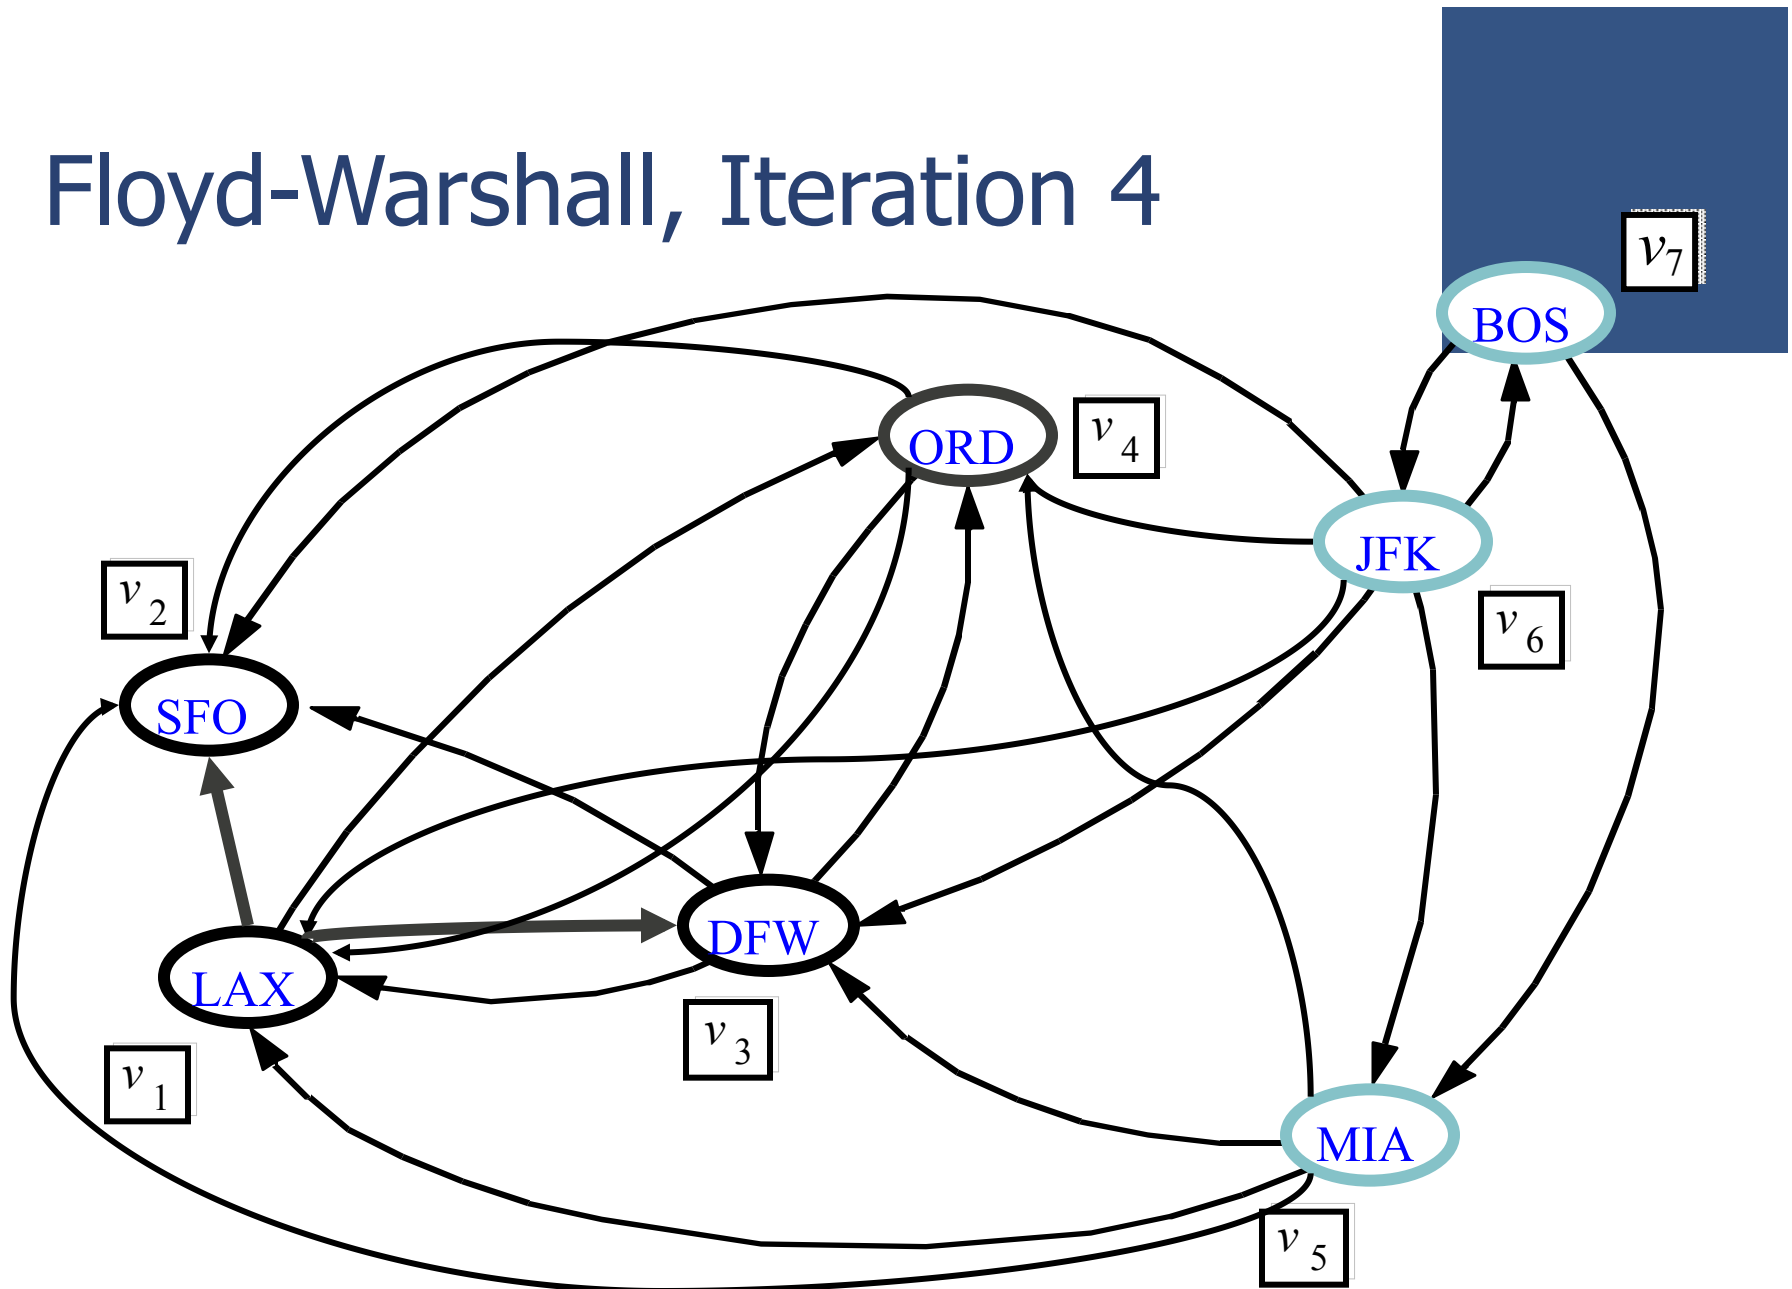
Floyd-Warshall, Iteration 2



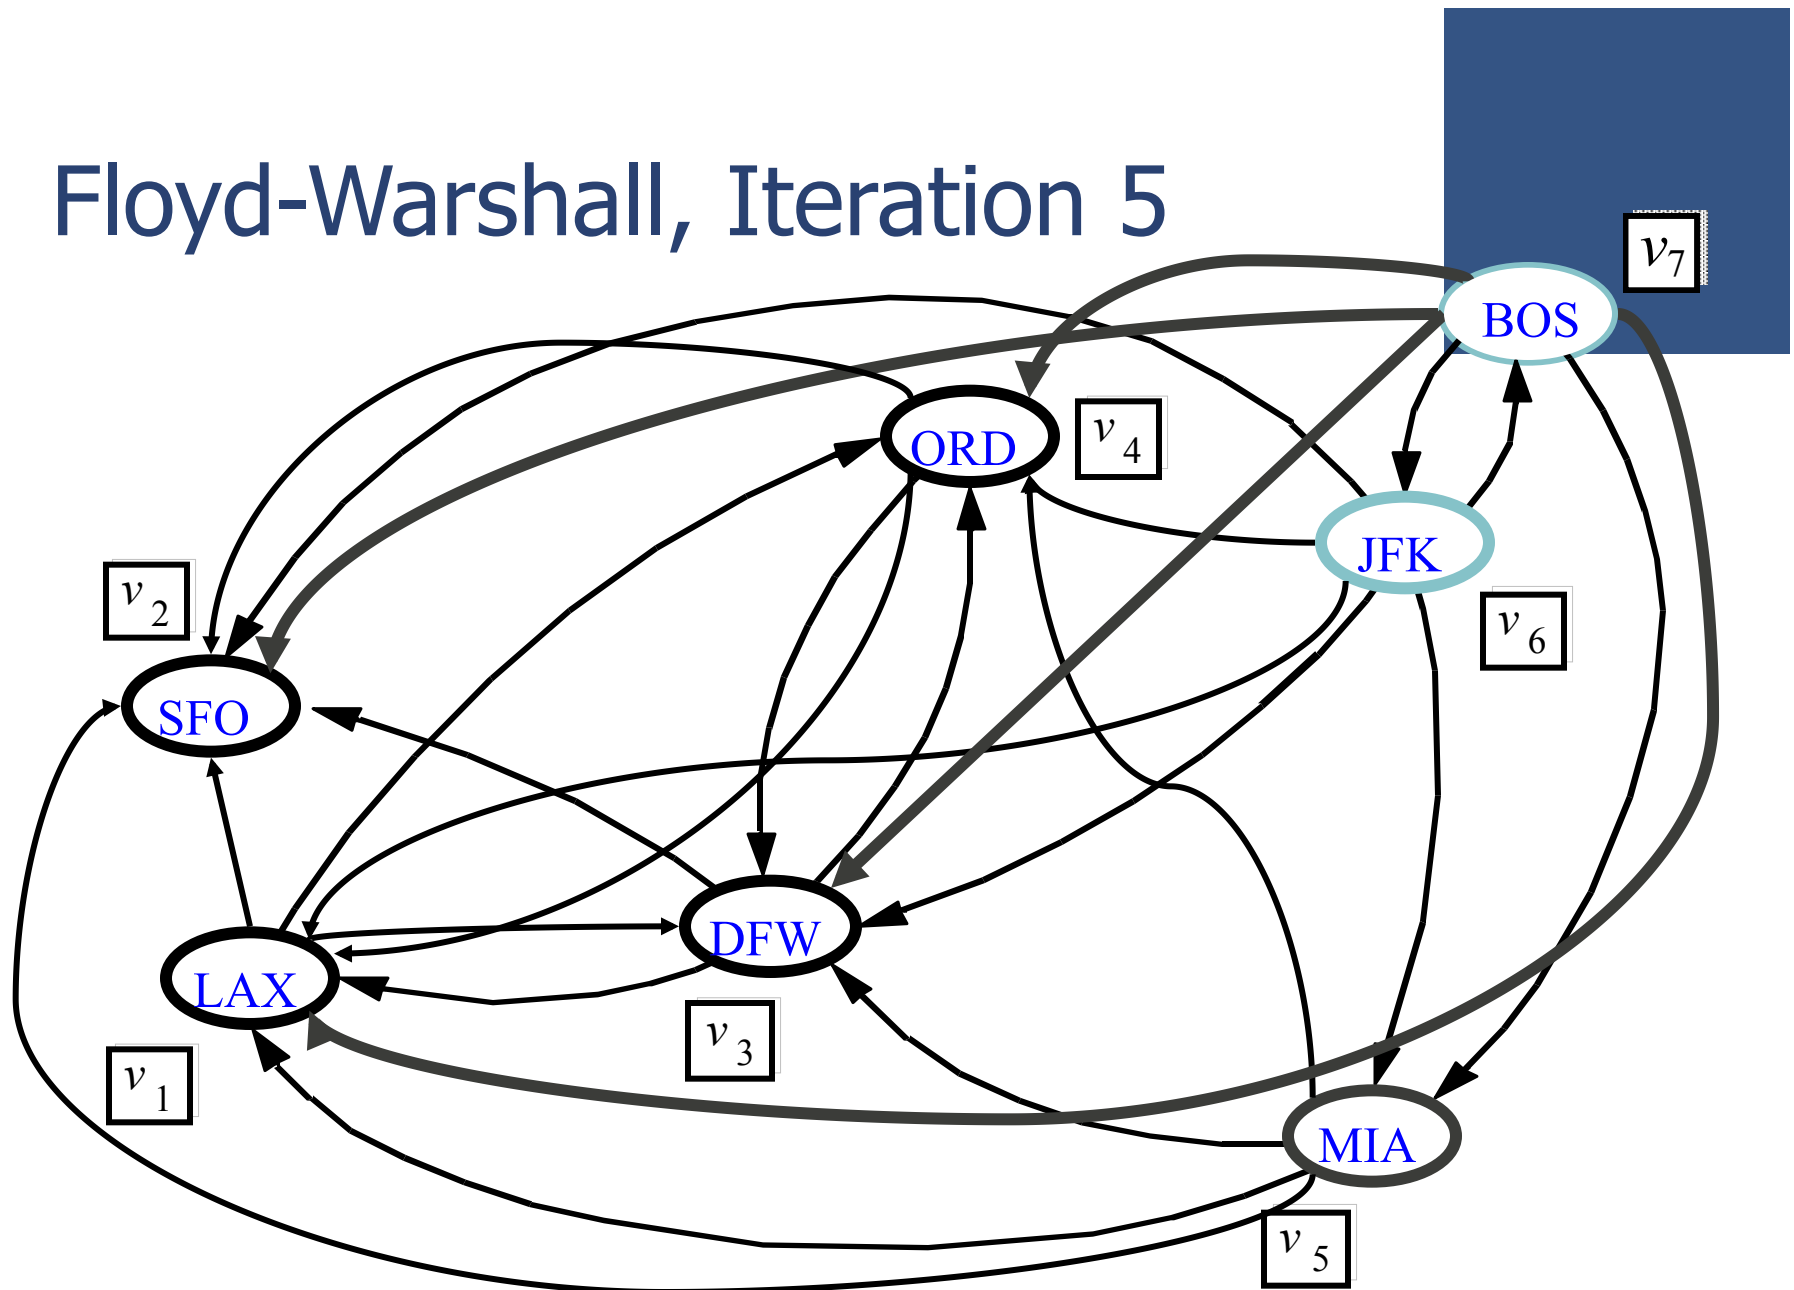
Floyd-Warshall, Iteration 3



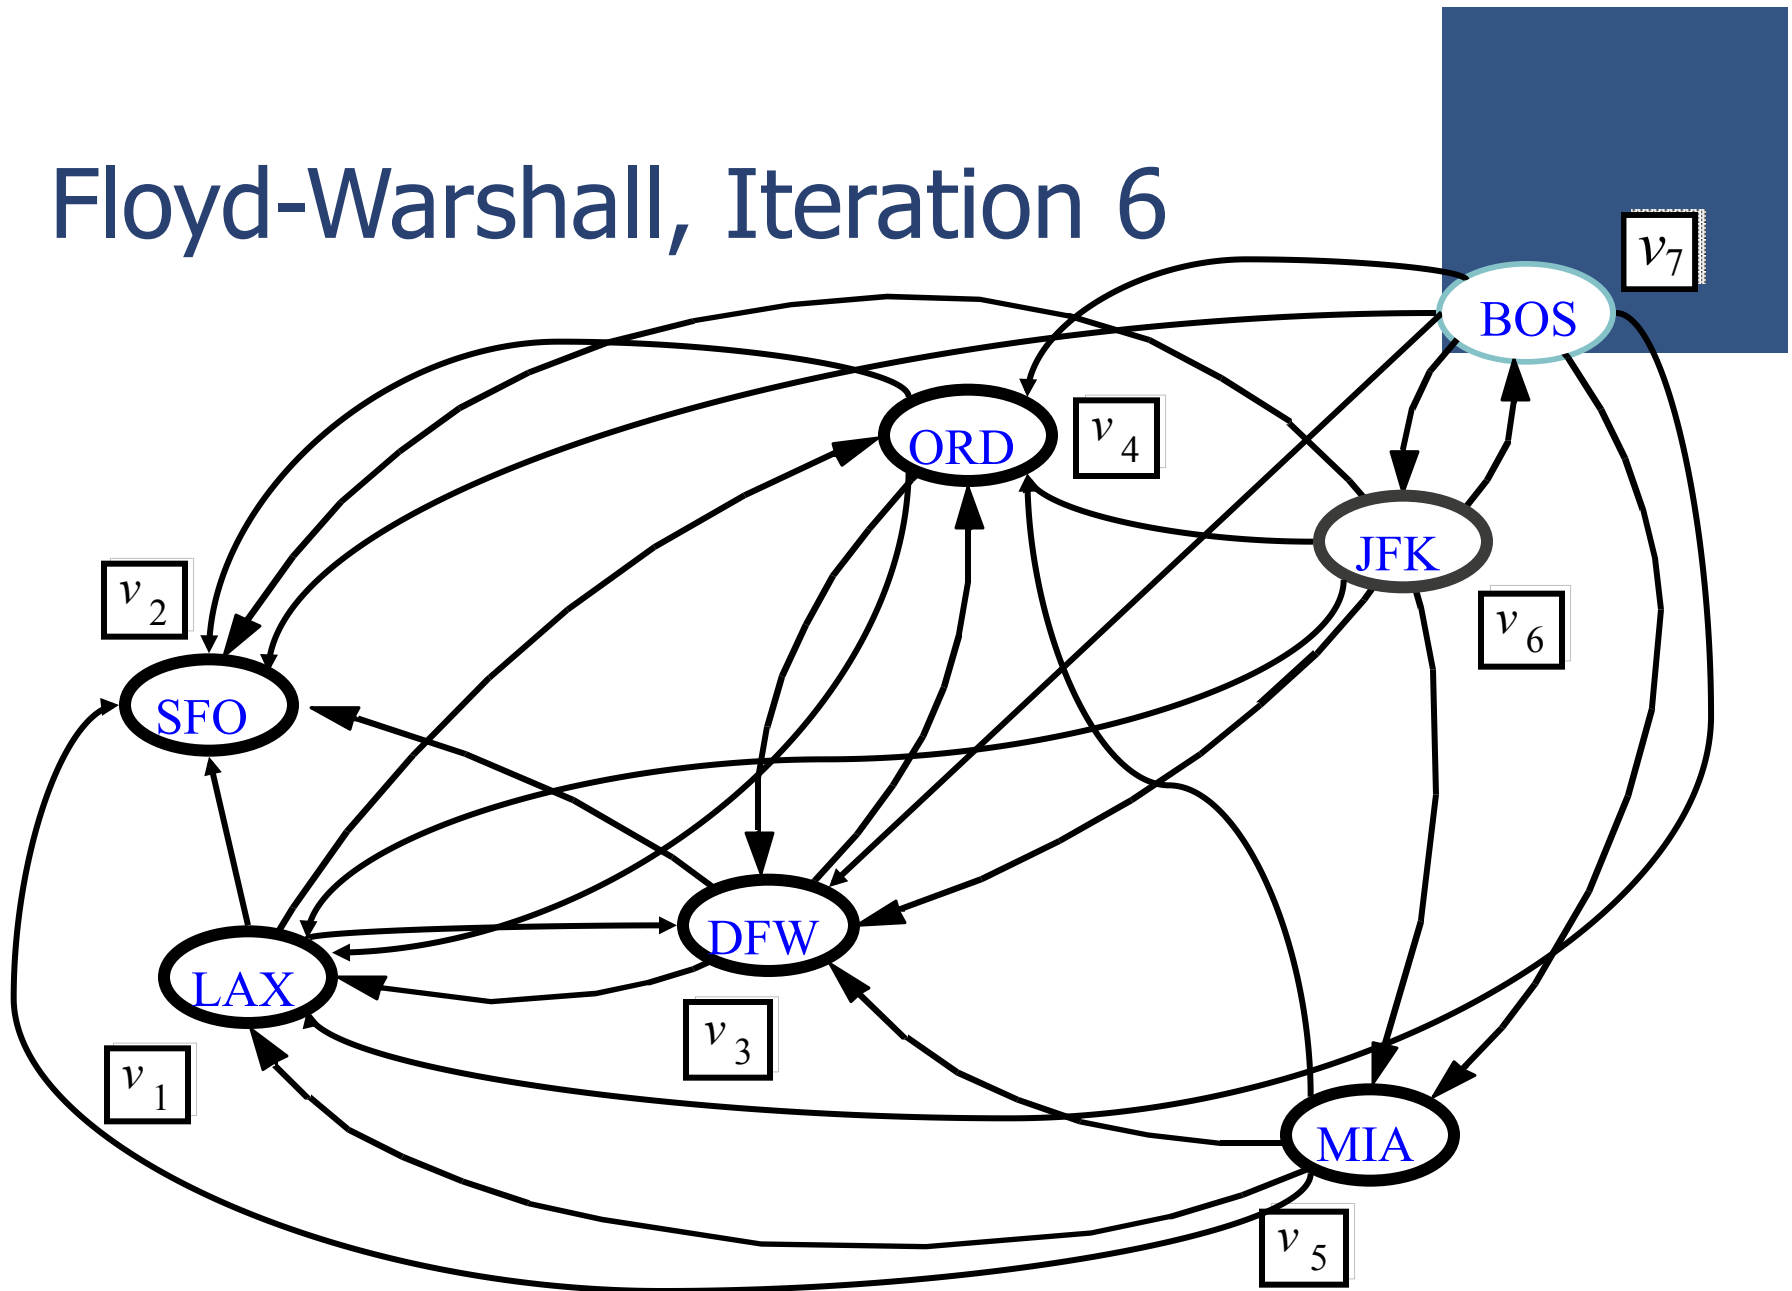
Floyd-Warshall, Iteration 4



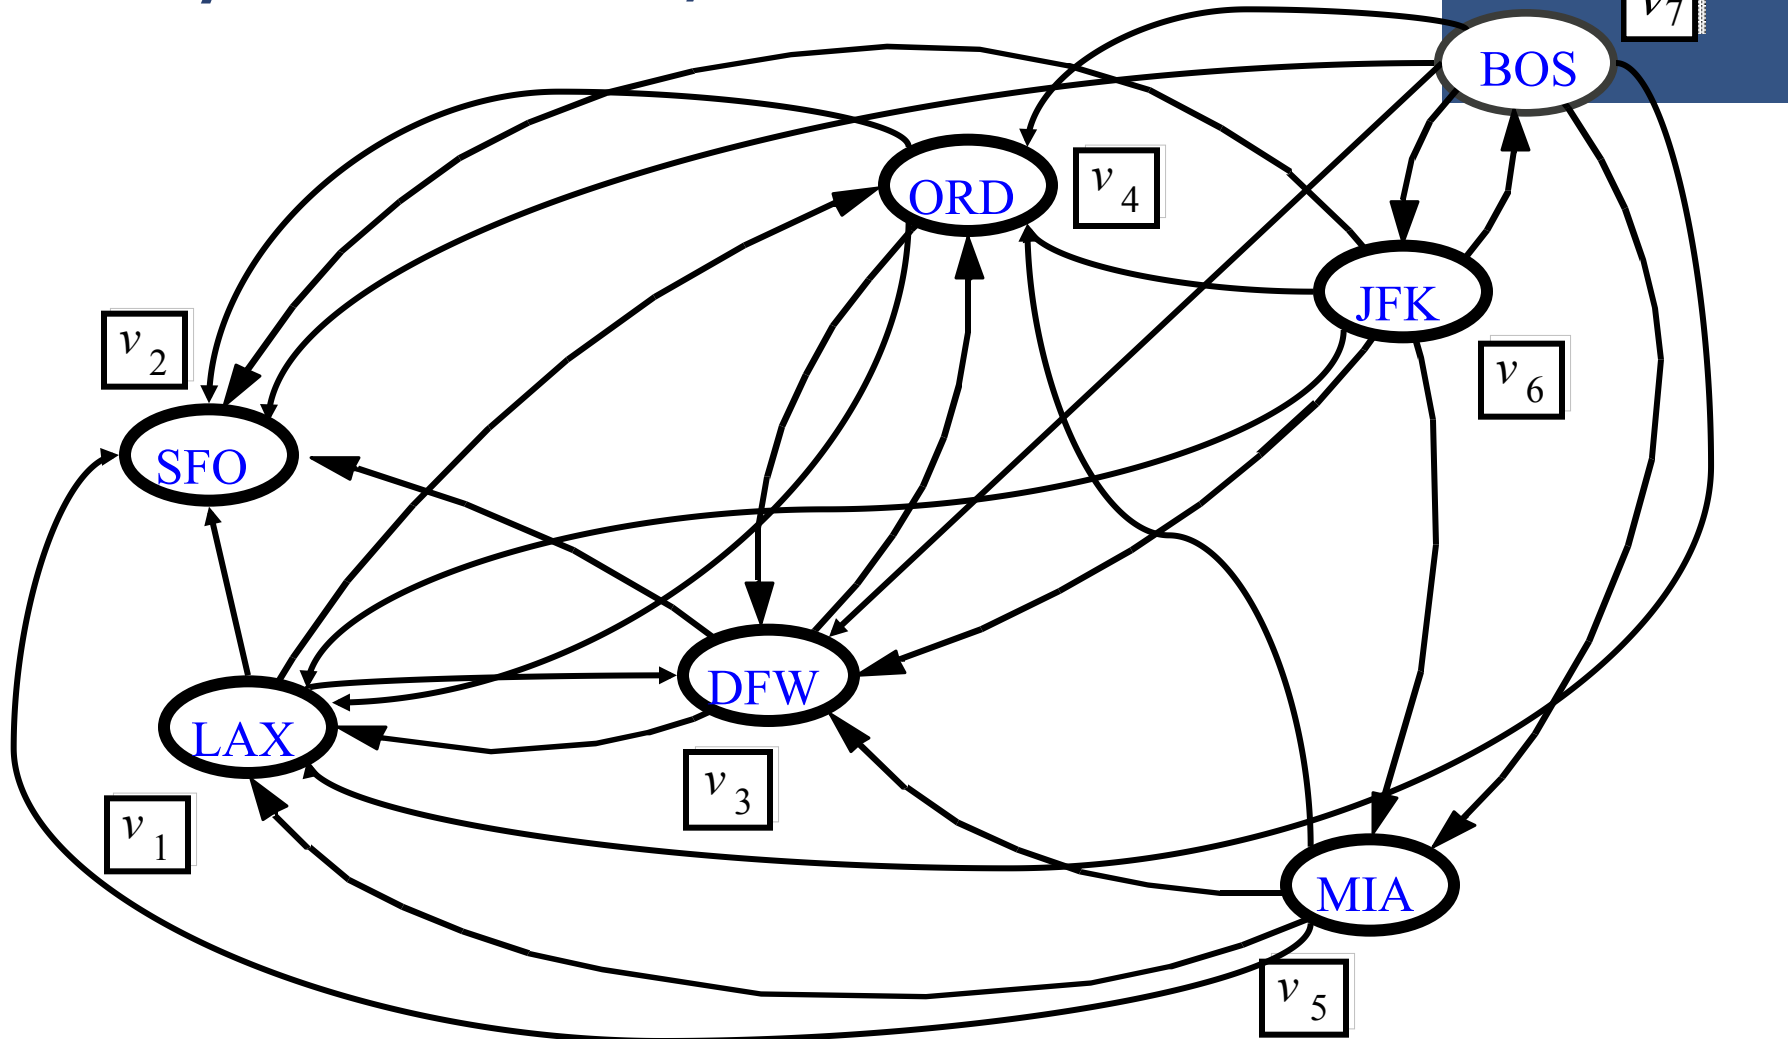
Floyd-Warshall, Iteration 5



Floyd-Warshall, Iteration 6




Floyd-Warshall, Conclusion



DAGs and Topological Ordering

- A directed acyclic graph (DAG) is a digraph that has no directed cycles




DAGs and Topological Ordering

- A topological ordering of a digraph is a numbering v_1, \dots, v_n of the vertices such that for every edge (v_i, v_j) , we have $i < j$
- Example: in a task scheduling digraph, a topological ordering a task sequence that satisfies the precedence constraints

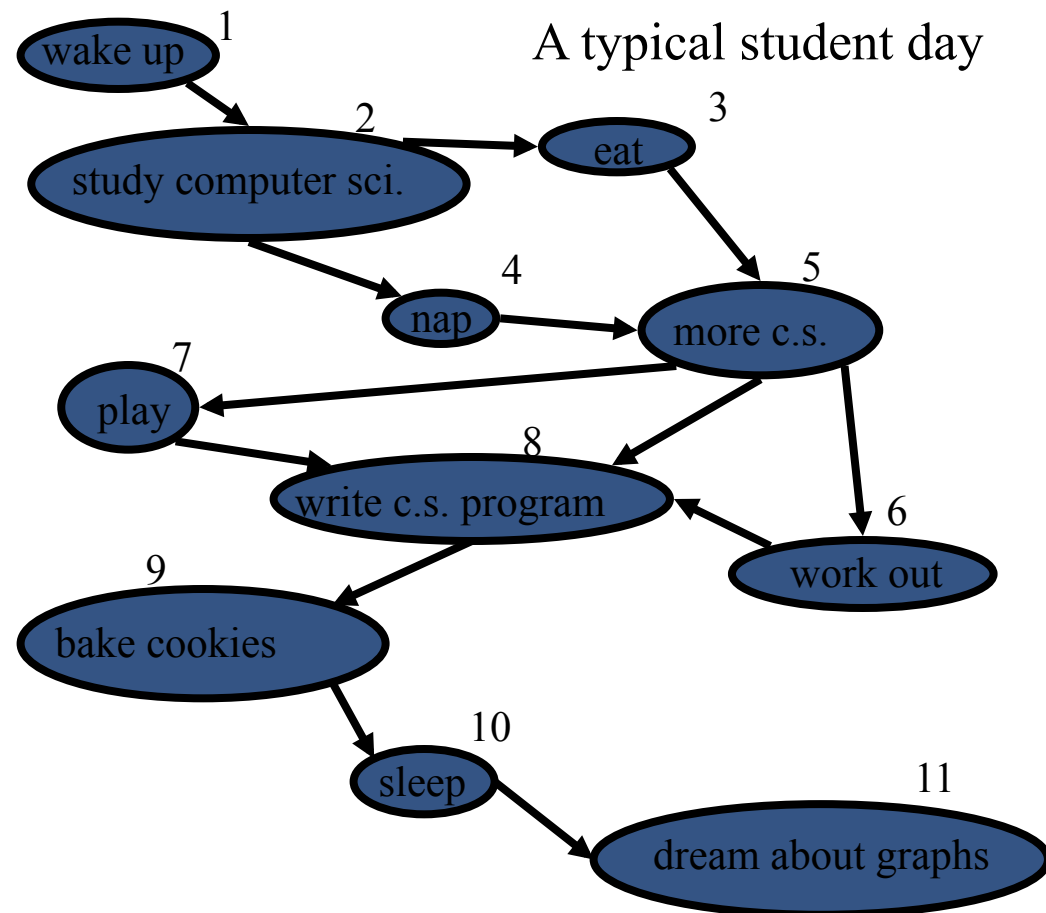
Theorem

A digraph admits a topological ordering if and only if it is a DAG



Topological Sorting

- Number vertices, so that (u,v) in E implies $u < v$



Algorithm for Topological Sorting

- Note: This algorithm is different than the one in the book

Algorithm TopologicalSort(G)

$H \leftarrow G$ // Temporary copy of G

$n \leftarrow G.\text{numVertices}()$

while H is not empty **do**

Let v be a vertex with no outgoing edges

Label $v \leftarrow n$

$n \leftarrow n - 1$

Remove v from H

- Running time: $O(n + m)$

Implementation with DFS

- Simulate the algorithm by using depth-first search
- $O(n+m)$ time.

Algorithm *topologicalDFS*(G)

Input dag G

Output topological ordering of G
 $n \leftarrow G.numVertices()$

for all $u \in G.vertices()$

setLabel(u , *UNEXPLORED*)

for all $v \in G.vertices()$

if *getLabel*(v) = *UNEXPLORED*
topologicalDFS(G , v)

Algorithm *topologicalDFS*(G , v)

Input graph G and a start vertex v of G

Output labeling of the vertices of G
in the connected component of v

setLabel(v , *VISITED*)

for all $e \in G.outEdges(v)$

{ outgoing edges }

$w \leftarrow opposite(v, e)$

if *getLabel*(w) = *UNEXPLORED*

{ e is a discovery edge }

topologicalDFS(G , w)


else

{ e is a forward or cross edge }


Label v with topological number n

$n \leftarrow n - 1$


Topological Sorting Example




Topological Sorting Example




Topological Sorting Example




Topological Sorting Example




Topological Sorting Example




Topological Sorting Example




Topological Sorting Example




Topological Sorting Example



Topological Sorting Example



Topological Sorting Example



A Quiz



- Liang wants to plan his course schedule. The course prerequisites are:

- CS15: (none)
- CS16: CS15
- CS22: (none)
- CS31: CS15
- CS32: CS16, CS31
- CS126: CS22, CS32, CS16
- CS127: CS16
- CS141: CS22, CS16
- CS169: CS32

Please help Liang to find the sequence of courses that allows him to satisfy all the prerequisites.

Minimum Spanning Trees

Spanning subgraph

- Subgraph of a graph G containing all the vertices of G

Spanning tree

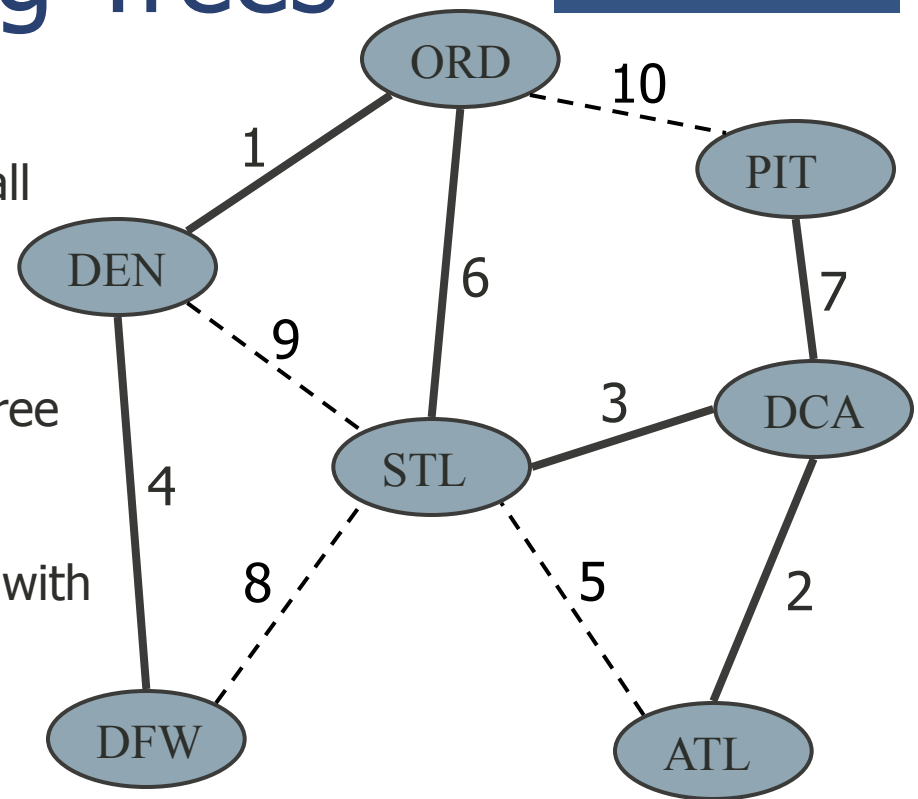
- Spanning subgraph that is itself a tree

Minimum spanning tree (MST)

- Spanning tree of a weighted graph with minimum total edge weight

■ Applications

- Communications networks
- Transportation networks




Cycle Property

Cycle Property:


- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T and C let be the cycle formed by e with T
- For every edge f of C , $weight(f) \leq weight(e)$

Proof:

- By contradiction
- If $weight(f) > weight(e)$ we can get a spanning tree of smaller weight by replacing e with f




Replacing f with e yields a better spanning tree




Partition Property

- Partition Property:
 - Consider a partition of the vertices of G into subsets U and V
 - Let e be an edge of minimum weight across the partition
 - There is a minimum spanning tree of G containing edge e



Replacing f with e yields another MST



Kruskal's Algorithm

- Maintain a partition of the vertices into clusters
 - Initially, single-vertex clusters
 - Keep an MST for each cluster
 - Merge “closest” clusters and their MSTs
- A priority queue stores the edges outside clusters
 - Key: weight
 - Element: edge
- At the end of the algorithm
 - One cluster and one MST

Algorithm *KruskalMST*(G)

for each vertex v in G **do**

 Create a cluster consisting of v

let Q be a priority queue.

Insert all edges into Q

$T \leftarrow \emptyset$

{ T is the union of the MSTs of the clusters}

while T has fewer than $n - 1$ edges **do**

$e \leftarrow Q.removeMin().getValue()$

$[u, v] \leftarrow G.endVertices(e)$

$A \leftarrow getCluster(u)$

$B \leftarrow getCluster(v)$



if $A \neq B$ **then**

 Add edge e to T




$mergeClusters(A, B)$

return T



Example




Example (contd.)



two steps



four steps



Prim-Jarnik's Algorithm

- We pick an arbitrary vertex s and we grow the MST as a cloud of vertices, starting from s
- We store with each vertex v label $d(v)$ representing the smallest weight of an edge connecting v to a vertex in the cloud
- At each step:
 - We add to the cloud the vertex u outside the cloud with the smallest distance label
 - We update the labels of the vertices adjacent to u


Prim-Jarnik's Algorithm (cont.)

- A heap-based adaptable priority queue with location-aware entries stores the vertices outside the cloud
 - Key: distance
 - Value: vertex
 - Recall that method *replaceKey* (l, k) changes the key of entry l
- We store three labels with each vertex:
 - Distance
 - Parent edge in MST
 - Entry in priority queue


Algorithm *PrimJarnikMST*(G)

```
 $Q \leftarrow$  new heap-based priority queue  
 $s \leftarrow$  a vertex of  $G$   
for all  $v \in G.vertices()$   
  if  $v = s$   
     $setDistance(v, 0)$   
  else  
     $setDistance(v, \infty)$   
     $setParent(v, \emptyset)$   
     $l \leftarrow Q.insert(getDistance(v), v)$   
     $setLocator(v, l)$   
while  $\neg Q.isEmpty()$   
   $l \leftarrow Q.removeMin()$   
   $u \leftarrow l.getValue()$   
  for all  $e \in G.incidentEdges(u)$   
     $z \leftarrow G.opposite(u, e)$   
     $r \leftarrow weight(e)$   
    if  $r < getDistance(z)$   
       $setDistance(z, r)$   
       $setParent(z, e)$   
       $Q.replaceKey(getEntry(z), r)$ 
```

Example



Example (contd.)



Baruvka's Algorithm

- Like Kruskal's Algorithm, Baruvka's algorithm grows many clusters at once and maintains a forest T
- Each iteration of the while loop halves the number of connected components in forest T
- The running time is $O(m \log n)$

Algorithm *BaruvkaMST*(G)

```
 $T \leftarrow V$  {just the vertices of  $G$ }  
while  $T$  has fewer than  $n - 1$  edges do  
  for each connected component  $C$  in  $T$  do  
    Let edge  $e$  be the smallest-weight edge from  $C$  to another component in  $T$   
    if  $e$  is not already in  $T$  then  
      Add edge  $e$  to  $T$   
return  $T$ 
```


Example of Baruvka's Algorithm (animated)

